



FRAUDOLENT TRANSACTION CLASSIFICATION

RICCARDO LA MARCA, MATR. 1795030

BIG DATA COMPUTING COURSE A.A. 2021/2022

1

Introduction

A brief presentation of the addressed problem

2

Dataset

A brief description of the dataset used in the project

3

Explore and Feature Engineering

How the dataset was modified

4

Machine Learning Models

The ML models and Pipelines applied for the task

5

Results

A description of the results obtained from the previous step

OVERVIEW



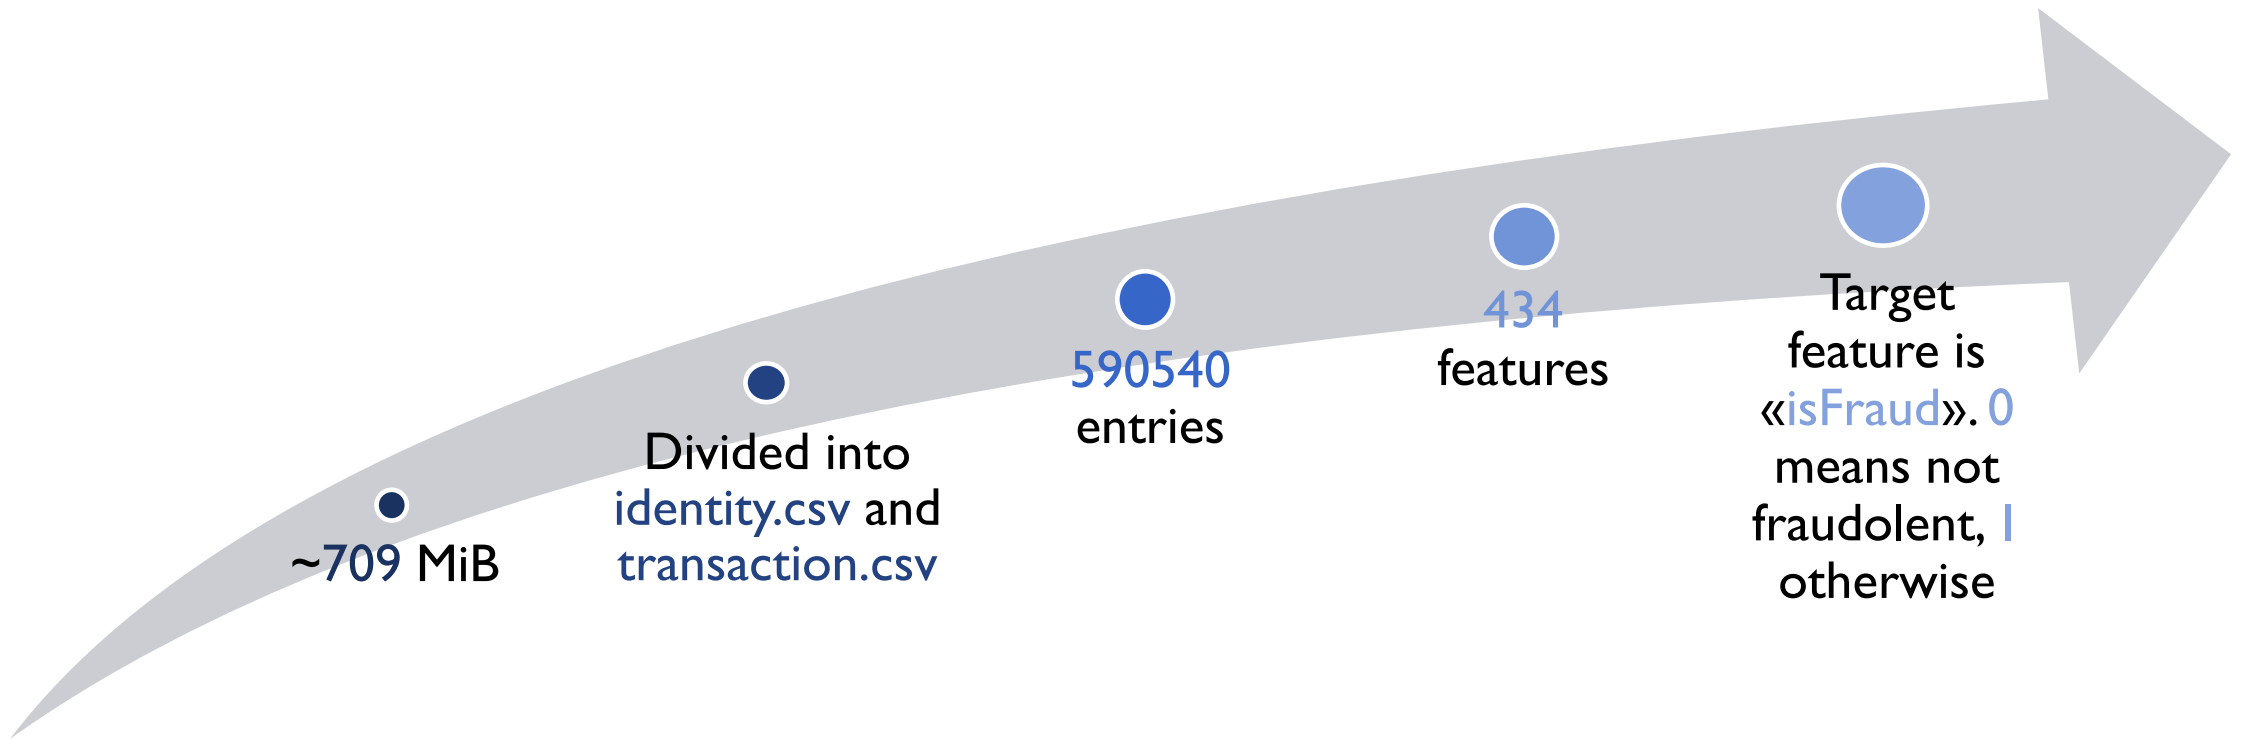
ADDRESSED PROBLEM

Financial fraud is a problem that has a huge impact on the financial industry

Credit card fraud detection is a challenge mainly due to 2 problems that it poses

- Both profiles of fraudulent and normal behaviours change
- Usually used datasets are highly skewed

The goal of the task is to create a Machine Learning model that, given a set of samples of fraudulent and not fraudulent transactions, is capable of classifying whether a new transaction is fraudulent or not.



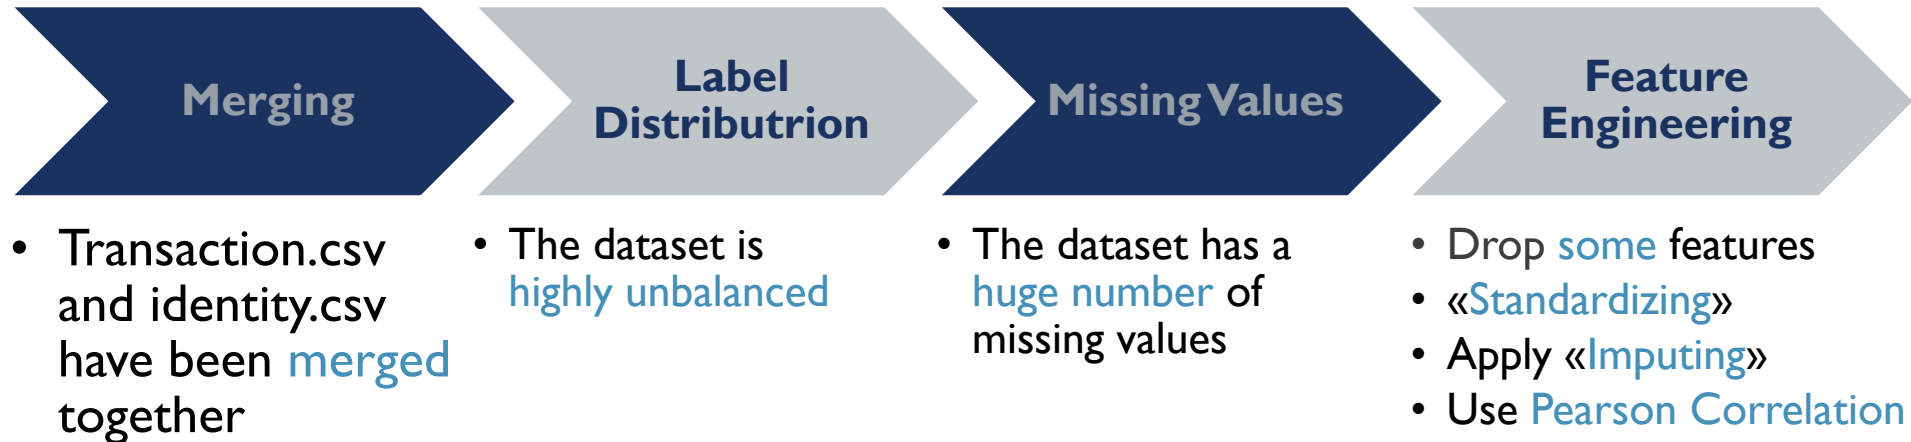
2

THE DATASET

The Dataset is available on [Kaggle](#)

3

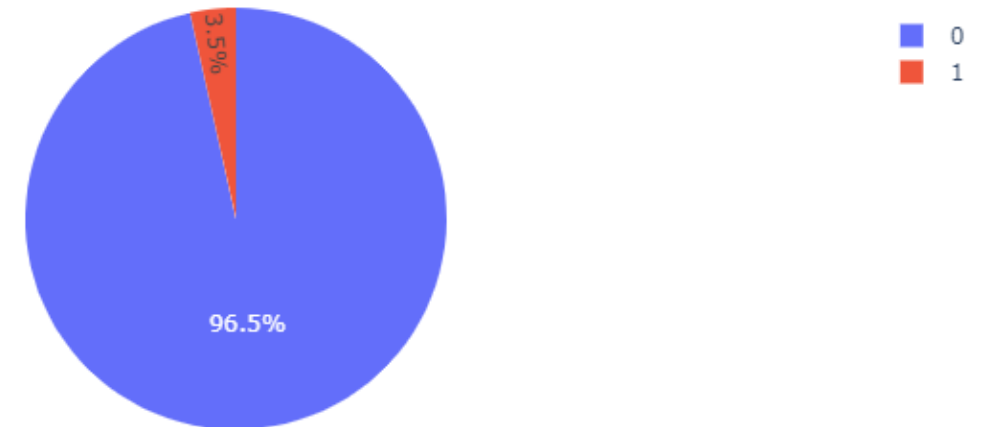
EXPLORE AND FEATURE ENGINEERING OUTLINE



3

.1 - LABEL DISTRIBUTION

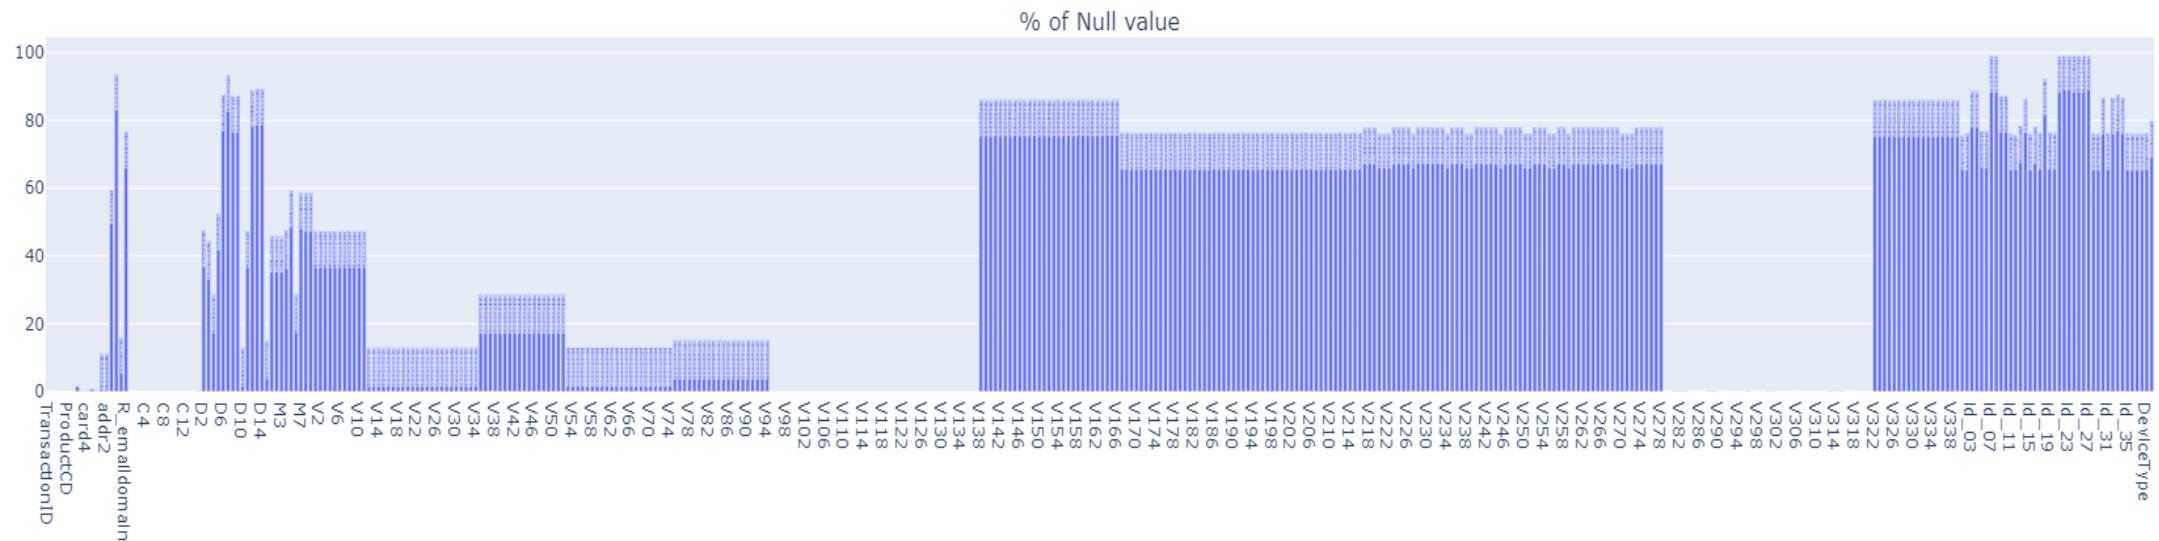
- With respect to the target label «isFraud» the dataset results **highly unbalanced**
- ~**96.5 %** are not-fraudulent transactions
- ~**3.5 %** are fraudulent transactions
- We have to handle this problem when splitting the dataset for training and testing the various ML models



3

.2 – MISSING VALUES

- The dataset has a **high number** of features with a huge percentage **of missing values**
- The average range of percentages is **~70-90%**
- I handled this during the Feature Engineering step



Features Dropping

- Drop features with percentage value of missing values greater or equal to 90%

Standardization

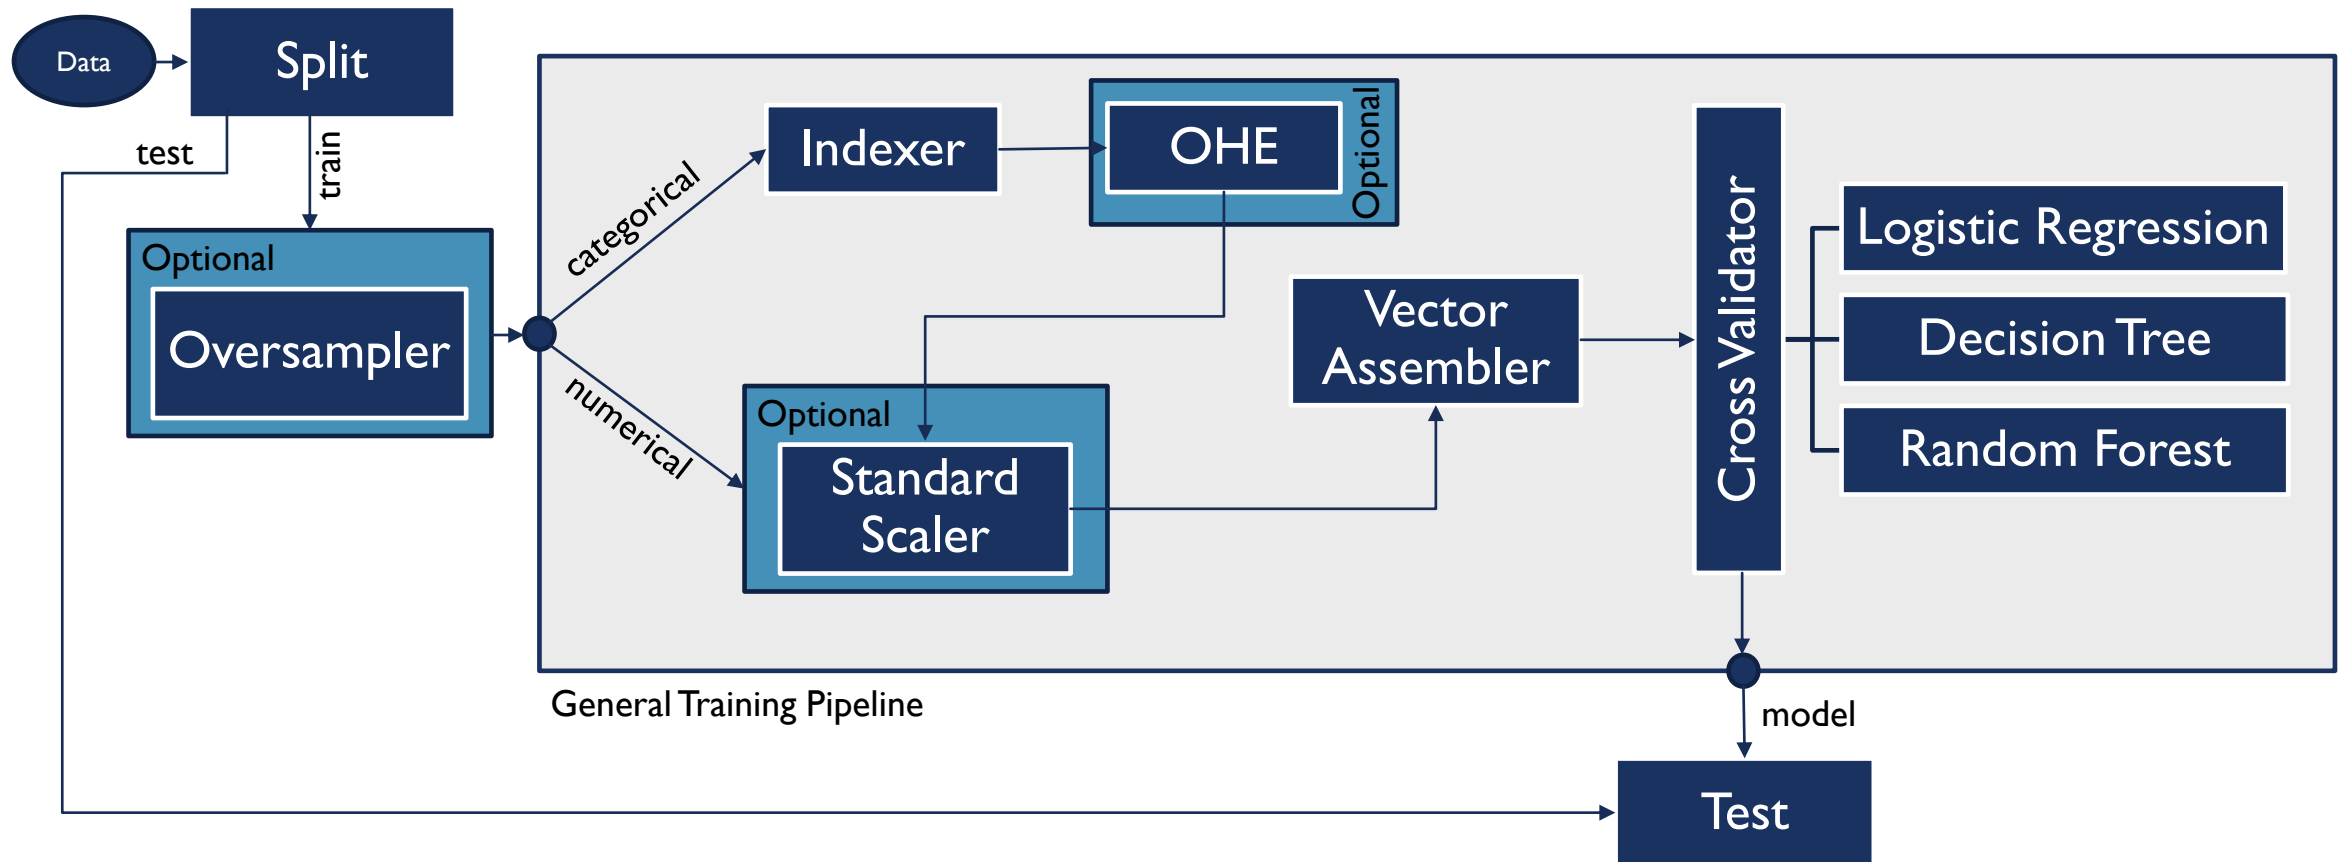
- Standardize certain features
- Given different values for the same feature but with equal meaning, replace with a single more general value
- Take *yahoo.co.jp*, *yahoo.co.uk* and *yahoo.net*, I replace it with *yahoo*

Imputing

- Use the imputer to replace null values in the dataset according to a specific strategy
- Discrete values use strategy *mean*
- Nulls in categorical values have been replaced with «N»

Pearson Correlation

- Drop more features using the Pearson Correlation
- If the PC > .95, then drop that feature

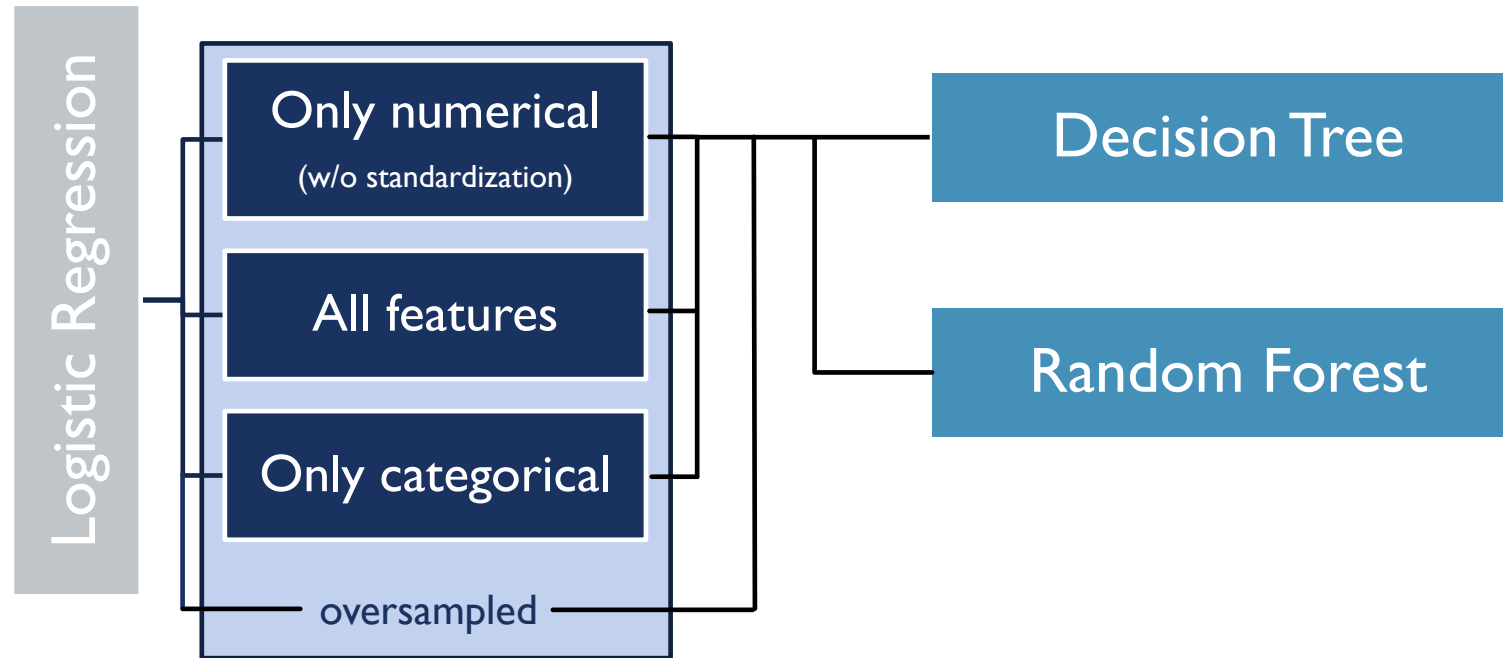


- The dataset is highly unbalanced, thus we cannot apply a **simple random splitting**
- This might lead to a **poor splitting strategy**
 - For instance the test set ends up containing only examples that are labeled with the most representative class
 - In this case such a class is the one for *non-fraudulent transactions*
- For this reason I used the so-called **Stratified Random Sampling**
 - It guarantees that both the training and the test split follow the same class distribution of the original dataset
 - For the experiments I selected **60%** of 0's and **70%** of 1's
- After splitting we last with: **357041 × 232** (train set) and **233499 × 232** (test set)

- After the stratified sampling, the train dataset was **still highly unbalanced**
- We have **342530** of 0s and **14511** of 1s.
- I decided to apply **oversampling** on the train set
 - After this I had 342530 of equal entries for fraudulent transactions
 - I decided to keep only the **60%** of them
 - That because, keeping all of them, I obtained a high number of False Positive
 - A high number of non fraudulent have been classified as fraudulent
- Finally, the train set contains **63.3%** of non fraud and **36.7%** of fraudulent

4

.2 – TRAIN OVERSAMPLER



Accuracy	Numerical (oversample)		All Features (oversample)	Categorical (oversample)
	with standardization	w/out standardize		
Logistic Regression	0.9772 (0.9099)	0.97725 (0.9099)	0.9777 (?)	0.9733 (0.8604)
Decision Tree	0.9773 (0.8586)	0.9773 (0.8586)	0.9790 (0.8586)	0.9734 (0.7862)
Random Forest	0.9787 (0.9390)	0.9789 (0.9397)	0.9786 (0.9410)	0.9734 (0.8829)

AUC ROC	Numerical (oversample)		All Features (oversample)	Categorical (oversample)
	with standardization	w/out standardize		
Logistic Regression	0.832 (0.840)	0.834 (0.840)	0.857 (?)	0.800 (0.801)
Decision Tree	0.428 (0.535)	0.428 (0.535)	0.858 (0.535)	0.707 (0.679)
Random Forest	0.844 (0.864)	0.845 (0.866)	0.848 (0.872)	0.784 (0.810)

F1-Score	Numerical (oversample)		All Features (oversample)	Categorical (oversample)
	with standardization	w/out standardize		
Logistic Regression	0.7137 (0.6446)	0.7139 (0.6446)	0.7227 (?)	0.5897 (0.6109)
Decision Tree	0.7138 (0.6101)	0.7138 (0.6101)	0.7472 (0.6101)	0.5961 (0.5999)
Random Forest	0.7406 (0.6787)	0.7438 (0.6787)	0.7390 (0.6842)	0.4932 (0.6165)

P/R	Numerical (oversample)		All Features (oversample)	Categorical (oversample)
	with standardization	w/out standardize		
Logistic Regression	0.853/0.613 (0.572/0.737)	0.852/0.614 (0.572/0.737)	0.858/0.624 (?)	0.706/0.505 (0.542/0.699)
Decision Tree	0.878/0.601 (0.541/0.698)	0.878/0.601 (0.541/0.698)	0.869/0.654 (0.541/0.698)	0.732/0.502 (0.529/0.692)
Random Forest	0.935/0.613 (0.611/0.762)	0.937/0.616 (0.612/0.761)	0.930/0.613 (0.616/0.768)	0.486/0.5 (0.550/0.701)